

AI-Powered Documentation Crawler & Q/A System

A comprehensive FastAPI-based RAG (Retrieval-Augmented Generation) system for intelligently crawling documentation websites and providing Q&A capabilities with advanced vector search and multi-domain support.

Table of Contents

-  Key Features
-  System Architecture
-  Quick Start
-  Complete API Reference
-  Detailed Project Structure
-  Storage Architecture Deep Dive
-  Comprehensive Usage Guide
-  Core Components Deep Dive
-  Monitoring and Debugging
-  Deployment
-  License
-  Contributing
-  Support

Key Features

Intelligent Web Crawling

- Async recursive crawling with domain restriction
- Rate limiting and retry logic with exponential backoff
- Content extraction using traflatura for clean text
- Comprehensive metadata preservation (titles, URLs, links)
- Respects robots.txt and implements polite crawling practices

Advanced Storage Architecture

- **Dual format persistence:** JSON (machine-readable) + YAML (human-readable)
- **Domain-based isolation:** Separate storage per domain with organized folder structure
- **Individual document files:** Each page stored as separate file for granular access
- **Session tracking:** Complete crawl metadata and statistics

Sophisticated Embedding Pipeline

- **Multi-model support:** Gemini (primary) with sentence-transformers fallback
- **Intelligent chunking:** Configurable chunk size with overlap for optimal retrieval
- **Batch processing:** Efficient embedding generation for large document sets
- **FAISS vector stores:** High-performance similarity search per domain

Advanced RAG System

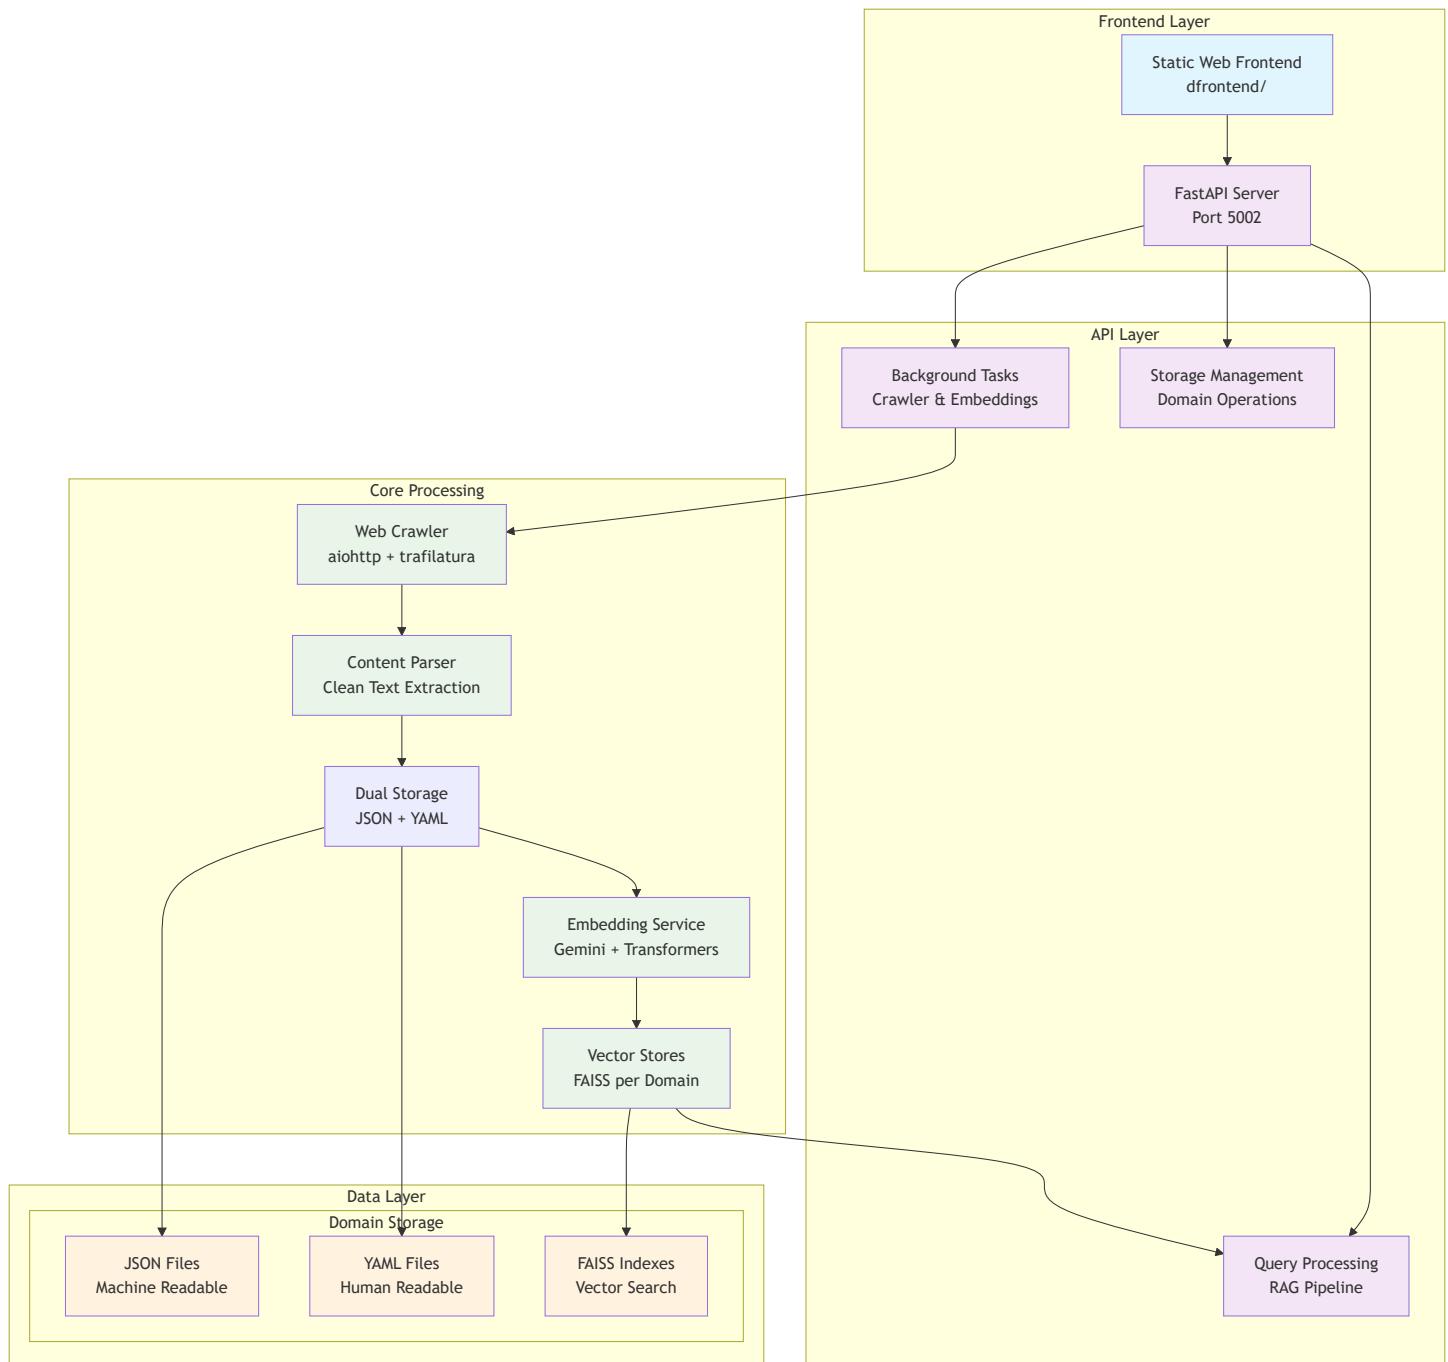
- **Multi-domain querying:** Search across multiple domains simultaneously
- **Context-aware generation:** Retrieval-augmented answers with source attribution
- **Configurable retrieval:** Adjustable top-k, context inclusion, and scoring
- **Background processing:** Async task management for long-running operations

Production-Ready API

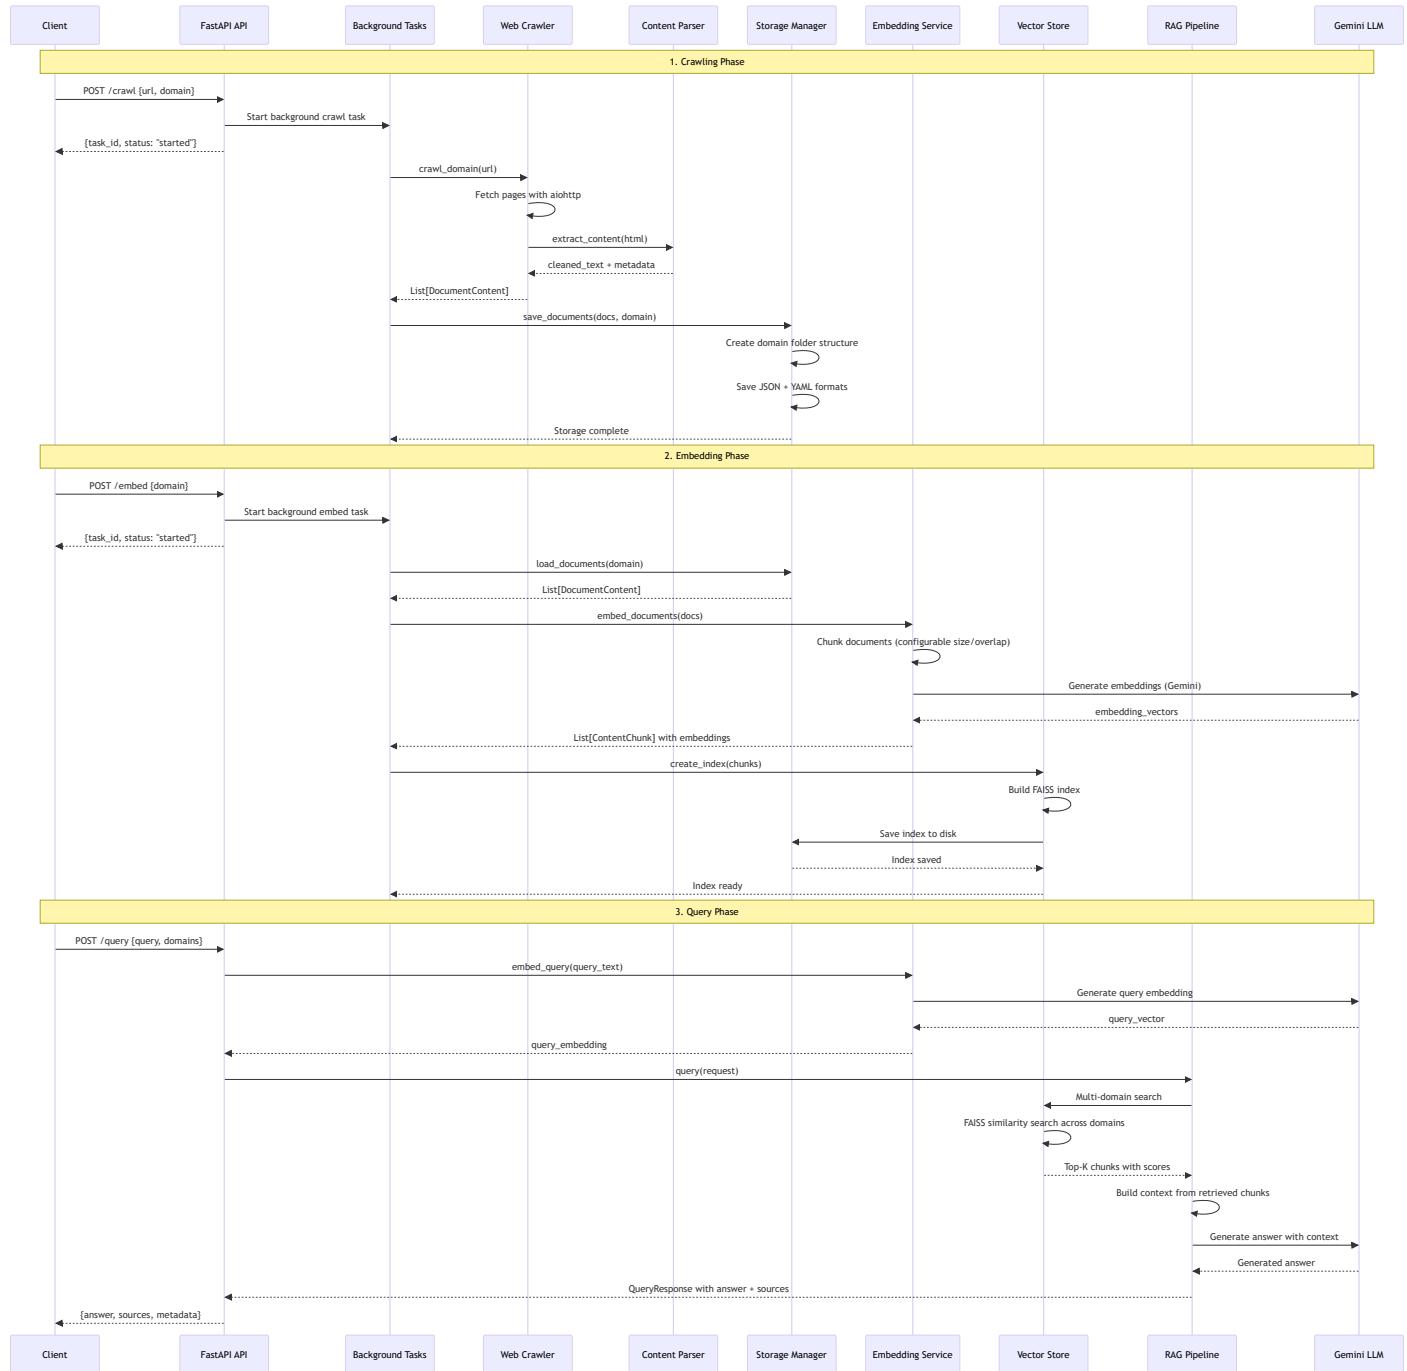
- **FastAPI framework:** Auto-generated OpenAPI docs and validation
- **Background task management:** Track crawling and embedding progress
- **CORS support:** Cross-origin requests for frontend integration
- **Comprehensive error handling:** Detailed error messages and status codes

System Architecture

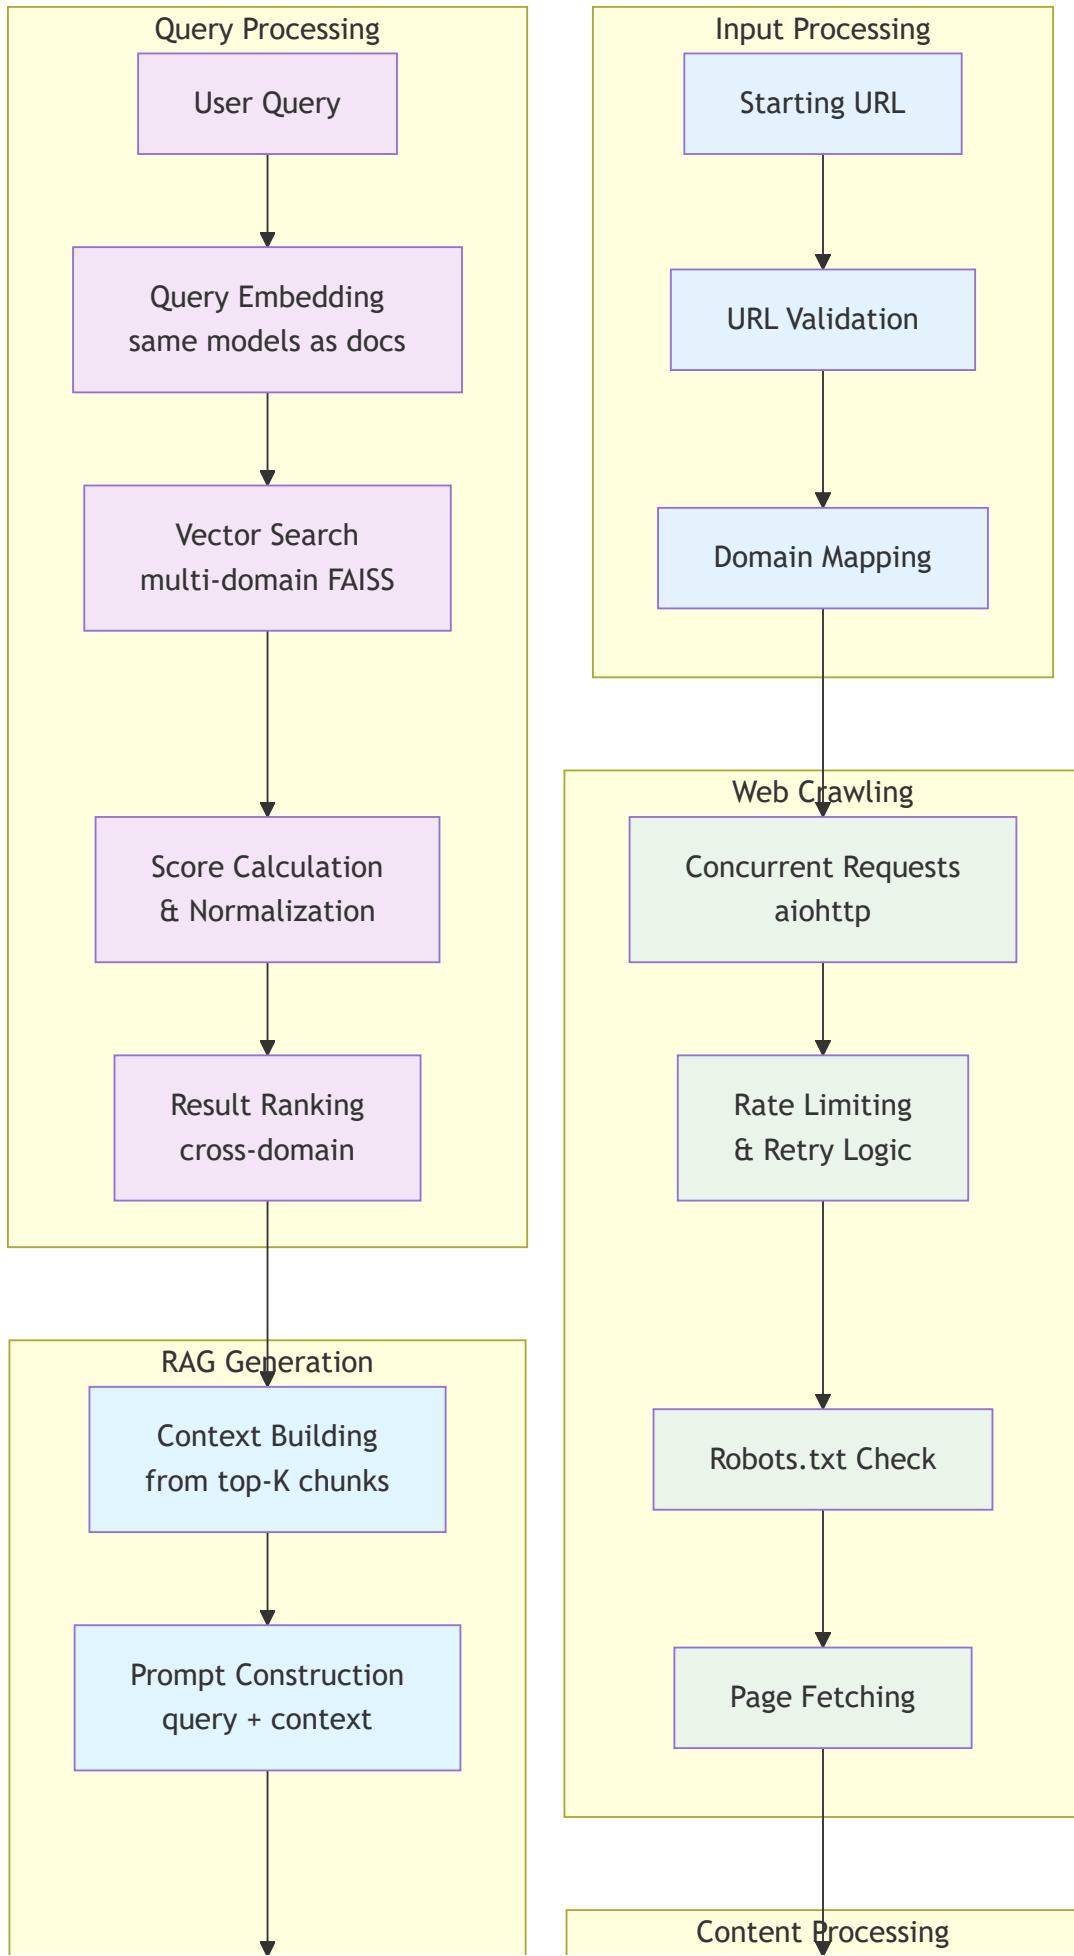
High-Level Architecture

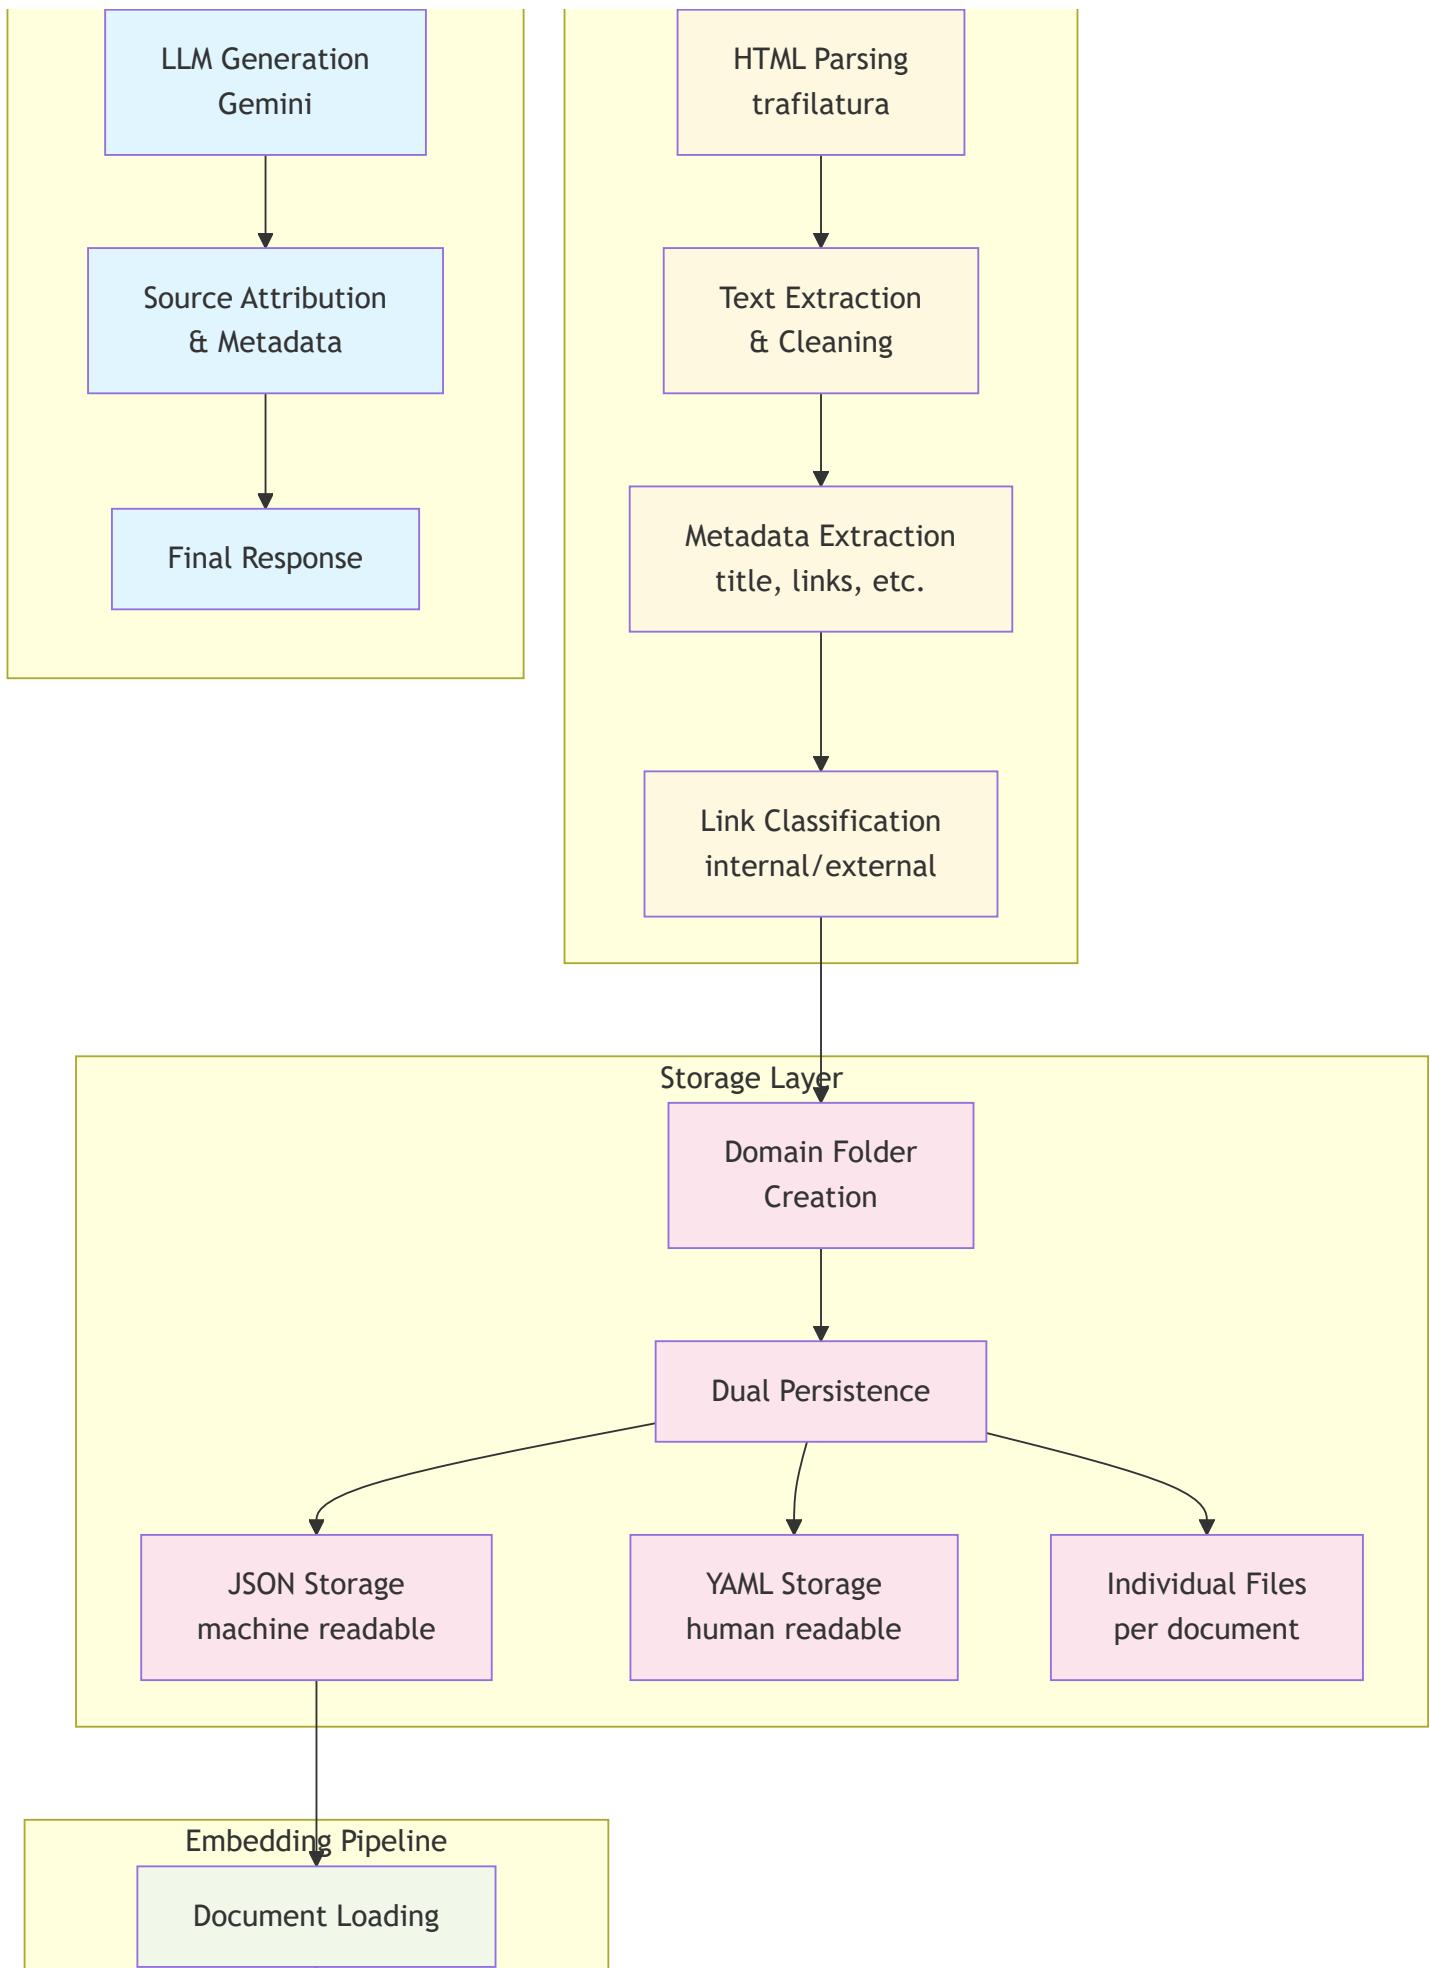


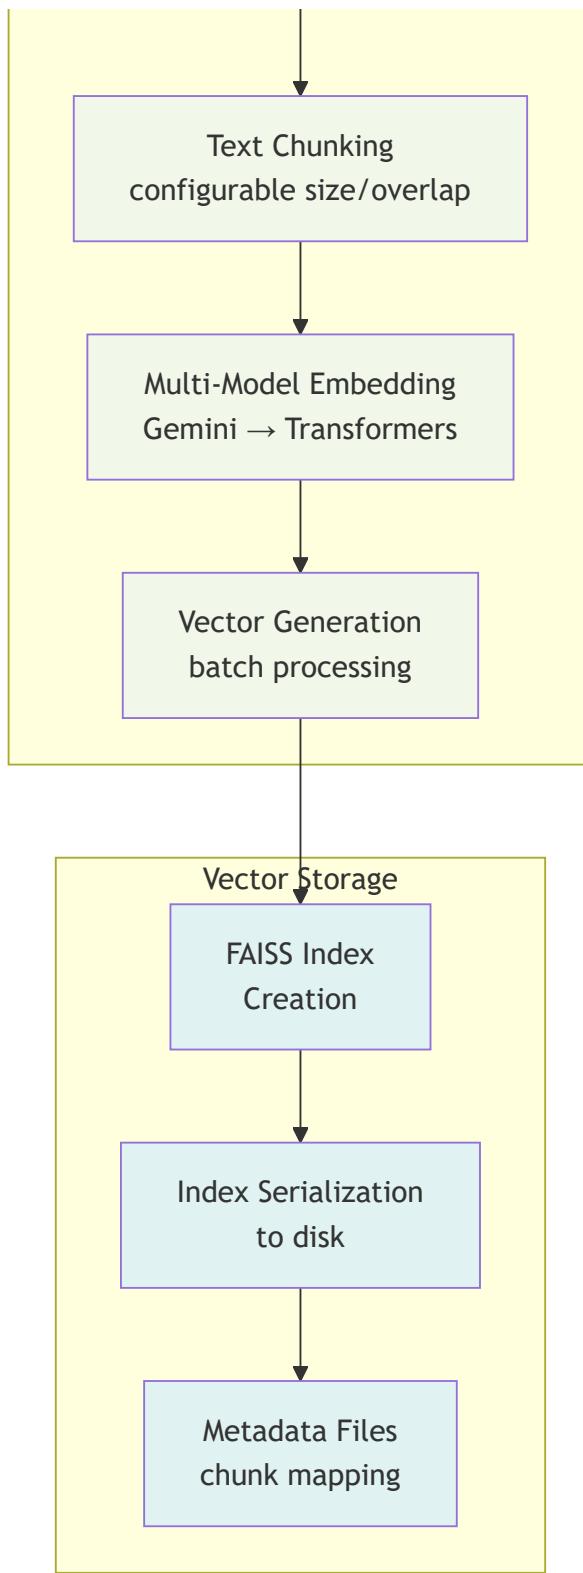
Detailed Code Flow



Data Flow Architecture







Quick Start

Prerequisites

- Python 3.8+

- Gemini API key from Google AI Studio
- Git for cloning the repository

1. Installation & Setup

```
# Clone the repository
git clone <your-repository-url>
cd docs-crawler

# Create virtual environment (recommended)
python -m venv venv

# Activate virtual environment
# On Windows:
venv\Scripts\activate
# On macOS/Linux:
source venv/bin/activate

# Install dependencies
pip install -r requirements.txt
```

2. Configuration

Create a `.env` file in the project root:

```
# Required: Gemini API Configuration
GEMINI_API_KEY=your_gemini_api_key_here

# Optional: Crawler Settings (defaults shown)
MAX_CONCURRENT_REQUESTS=10
REQUEST_TIMEOUT=30
RETRY_ATTEMPTS=3
DELAY_BETWEEN_REQUESTS=1

# Optional: Storage Settings
DATA_DIR=./data
LOG_LEVEL=INFO

# Optional: Embedding Settings
CHUNK_SIZE=1000
CHUNK_OVERLAP=100
MAX_CHUNKS_PER_DOC=50

# Optional: FastAPI Settings
API_HOST=0.0.0.0
API_PORT=5002
```

3. Start the System

Option A: Using the run server script (Recommended)

```
python run_server.py
```

Option B: Direct FastAPI execution

```
cd src
python main.py
```

Option C: Using the deployment script

```
# For full deployment with frontend
.\deploy.bat
```

The API will be available at:

- **API Server:** <http://localhost:5002>
- **API Documentation:** <http://localhost:5002/docs>
- **Frontend** (if deployed): <http://localhost:3000>

Complete API Reference

Endpoint Quick Reference

Area	Method	Path	Description
Health	GET	/	Basic health check
System Status	GET	/status	Overall system status and metrics
Start Crawl	POST	/crawl	Start asynchronous crawl for a domain
Task Status	GET	/tasks/{task_id}	Check status/progress of background tasks
Generate Embeddings	POST	/embed	Build/refresh embeddings for a domain
Single-Domain Query	POST	/query	Query a single domain
Multi-Domain Query	POST	/query/multi-domain	Query across multiple domains
List Domains	GET	/domains	List available domains and their status
Domain Documents	GET	/domains/{domain}/documents	List stored documents for a domain
Validate Domains	POST	/domains/validate	Validate domains for querying/embeddings
Domain Stats	GET	/domains/{domain}/stats	Detailed statistics for a specific domain

Area	Method	Path	Description
Bulk Crawl	POST	/bulk/crawl	Start multiple domain crawls in one request

System Endpoints

Health Check

GET /

Response: System status and available features

System Status

GET /status

Response: Detailed system information including active domains, task counts, and resource usage

Crawling Operations

Start Domain Crawling

POST /crawl

Content-Type: application/json

```
{
  "url": "https://docs.livekit.io",
  "domain_name": "livekit-docs",      # optional, auto-generated from URL if omitted
  "max_depth": 10,                  # optional, default: 5
  "max_pages": 1000                # optional, default: 500
}
```

Response: Task ID for tracking progress

Track Crawling Progress

GET /tasks/{task_id}

Response: Task status, progress, and completion details

Embedding Operations

Generate Domain Embeddings

```
POST /embed
Content-Type: application/json

{
  "domain": "livekit-docs",
  "force_rebuild": false,           # optional, default: false
  "chunk_size": 1000,              # optional, uses config default
  "chunk_overlap": 100            # optional, uses config default
}
```

Response: Task ID for embedding generation tracking

Query Operations

Single Domain Query

```
POST /query
Content-Type: application/json

{
  "query": "How do I set up authentication?",
  "domain": "livekit-docs",
  "top_k": 5,                      # optional, default: 5
  "include_context": true,          # optional, default: true
  "min_relevance_score": 0.7       # optional, default: 0.0
}
```

Multi-Domain Query

```
POST /query/multi-domain
Content-Type: application/json

{
  "query": "How to implement real-time features?",
  "domains": ["livekit-docs", "docs-streamlit-io"],
  "top_k": 8,                      # optional, total results across domains
  "per_domain_k": 4,                # optional, candidates per domain
  "include_context": true,
  "merge_strategy": "score_based"  # optional: "score_based" or "round_robin"
}
```

Domain Management

List Available Domains

```
GET /domains
```

Response: List of domains with crawling and embedding status

Get Domain Documents

```
GET /domains/{domain_name}/documents?format_type=json&limit=100&offset=0
```

Parameters:

- `format_type` : "json" or "yaml" (default: json)
- `limit` : Number of documents to return (default: 100)
- `offset` : Pagination offset (default: 0)

Validate Domains for Querying

```
POST /domains/validate
Content-Type: application/json
```

```
{
  "domains": ["livekit-docs", "streamlit-docs"]
}
```

Response: Validation status per domain with embedding availability

Advanced Operations

Domain Statistics

```
GET /domains/{domain_name}/stats
```

Response: Comprehensive domain statistics including document count, embedding status, index size, and performance metrics

Bulk Operations

```
POST /bulk/crawl
```

Content-Type: application/json

```
{
  "urls": [
    {"url": "https://docs.livekit.io", "domain_name": "livekit"},
    {"url": "https://docs.streamlit.io", "domain_name": "streamlit"}
  ],
  "max_concurrent": 2           # optional, default: 3
}
```

Detailed Project Structure

```
docs-crawler/
├── README.md                      # This comprehensive documentation
├── requirements.txt                 # Python dependencies
├── run_server.py                   # Main server launcher with path setup
├── deploy.bat                      # Full deployment script (backend + frontend)
└── .env.example                     # Environment variables template

|
├── src/                            # Core application source code
│   ├── main.py                     # FastAPI application entry point
│   ├── asgi.py                     # ASGI production configuration
│   └── wsgi.py                      # WSGI production configuration

|
│   ├── config/                     # Configuration management
│   │   ├── __init__.py
│   │   └── settings.py             # Environment-based settings with validation

|
│   ├── api/                         # FastAPI routes and models
│   │   ├── __init__.py
│   │   ├── endpoints.py            # All API route definitions and handlers
│   │   └── models.py               # Pydantic request/response models

|
│   ├── crawler/                    # Web crawling engine
│   │   ├── __init__.py
│   │   ├── web_crawler.py          # Async web crawler with rate limiting
│   │   └── content_parser.py       # Content extraction with trafiletura

|
│   ├── storage/                    # Data persistence layer
│   │   ├── __init__.py
│   │   ├── storage_manager.py      # Domain-based storage orchestration
│   │   ├── schemas.py              # Data schemas and models
│   │   └── azure_blob.py           # Cloud storage integration (future)

|
│   ├── embeddings/                 # Vector embeddings and search
│   │   ├── __init__.py
│   │   ├── embedding_service.py    # Multi-model embedding generation
│   │   ├── vector_store.py         # Single domain FAISS operations
│   │   └── multi_domain_vector_store.py # Cross-domain search orchestration

|
└── qa/                            # Question answering system
    └── __init__.py
```

```
|   |   └── rag_pipeline.py      # RAG implementation with context management
|   |
|   └── utils/                  # Utility modules
|       ├── __init__.py
|       ├── logger.py          # Structured logging configuration
|       └── url_utils.py       # URL normalization and validation
|
└── data/                      # Domain-organized data storage
    ├── {domain-name}/        # Auto-created per crawled domain
    |   ├── json/              # Machine-readable JSON files
    |   |   ├── crawl_session.json    # Crawl metadata and stats
    |   |   ├── documents.json      # All documents collection
    |   |   └── individual/        # Individual document files
    |   |       ├── doc-1.json
    |   |       └── doc-2.json
    |   ├── yaml/              # Human-readable YAML files
    |   |   ├── crawl_session.yaml
    |   |   ├── documents.yaml
    |   |   └── individual/
    |   └── faiss/              # Vector index files
    |       ├── index.faiss        # FAISS vector index
    |       ├── metadata.json      # Chunk-to-document mapping
    |       └── index_info.json    # Index configuration
    └── logs/                    # Application logs
        ├── crawler_YYYYMMDD.log
        ├── embedding_YYYYMMDD.log
        └── api_YYYYMMDD.log
|
└── dfrontend/                 # Static web frontend
    ├── index.html            # Main frontend application
    ├── package.json           # Node.js dependencies (for dev tools)
    ├── deploy.bat             # Frontend deployment script
    ├── run.bat                # Local frontend server
    ├── css/
    |   └── styles.css          # Application styles with dark/light themes
    ├── js/
    |   ├── config.js           # API endpoint configuration
    |   ├── main.js              # Core application logic
    |   ├── api.js               # API communication layer
    |   └── ui.js                # UI management and interactions
    └── assets/                 # Static assets
|
└── docs/                      # Project documentation
```

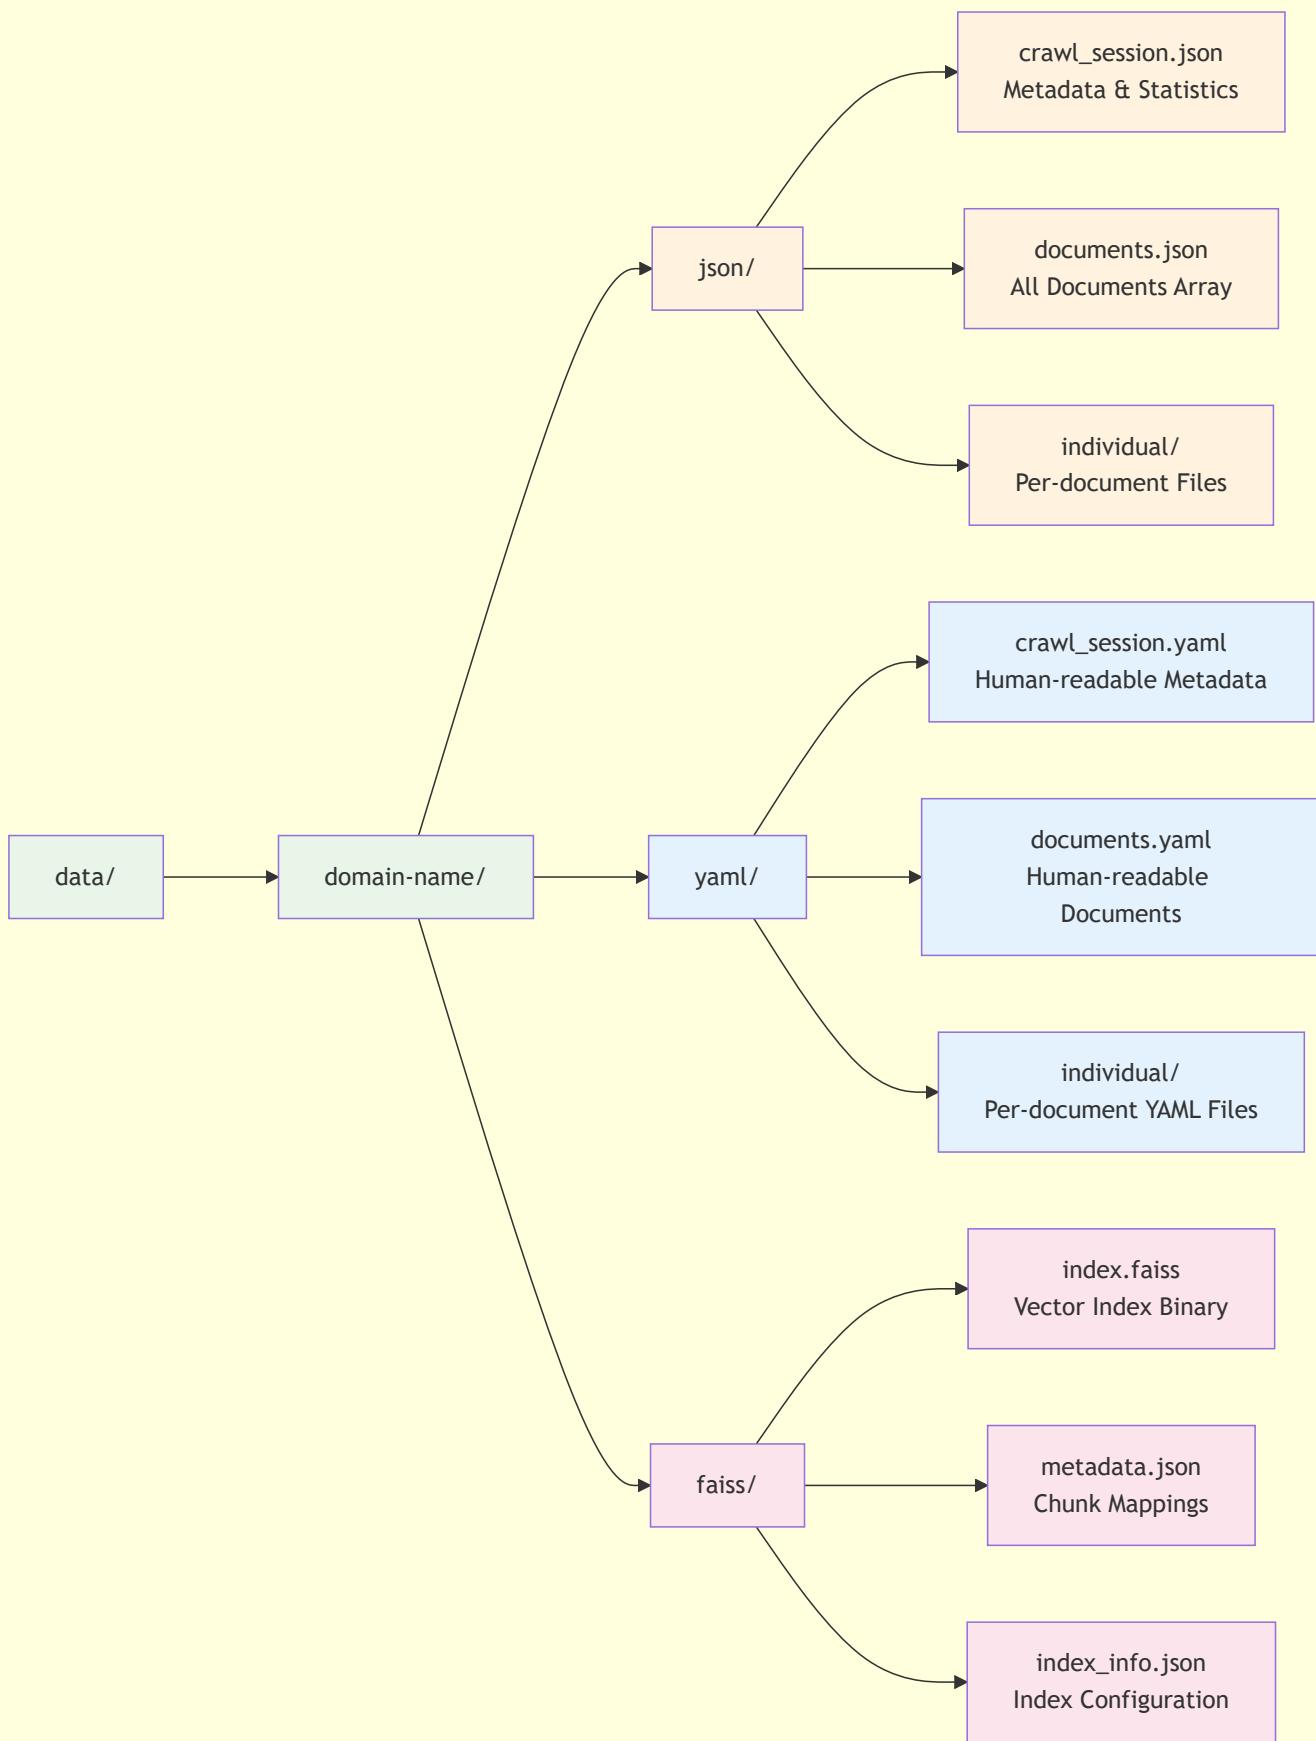
```
|   └── project_structure.md      # Detailed architecture documentation
|   └── AZURE_DEPLOYMENT.md      # Cloud deployment guide
|   └── *.md                      # Additional design documents
|
└── tests/                         # Test suite
    ├── test_crawl.py            # Crawling functionality tests
    ├── test_imports.py          # Import validation tests
    └── test_task_formatting.py  # Background task tests
```

Storage Architecture Deep Dive

Domain Folder Organization

Each crawled domain gets its own isolated storage structure:

Data Directory Structure



Storage Formats

1. JSON Format (Machine Processing)

```
{  
    "id": "doc-123",  
    "url": "https://docs.example.com/guide",  
    "title": "Getting Started Guide",  
    "content": "Clean extracted text...",  
    "metadata": {  
        "crawl_timestamp": "2024-01-15T10:30:00Z",  
        "content_length": 2847,  
        "internal_links": ["https://docs.example.com/api"],  
        "external_links": ["https://github.com/example/repo"]  
    },  
    "processing_info": {  
        "extraction_method": "trafilatura",  
        "chunk_count": 3,  
        "embedded": true  
    }  
}
```

2. YAML Format (Human Readable)

```
id: doc-123
url: https://docs.example.com/guide
title: Getting Started Guide
content: |
  Clean extracted text...

metadata:
  crawl_timestamp: 2024-01-15T10:30:00Z
  content_length: 2847
  internal_links:
    - https://docs.example.com/api
  external_links:
    - https://github.com/example/repo

processing_info:
  extraction_method: trafilatura
  chunk_count: 3
  embedded: true
```

3. FAISS Index Structure

```
{
  "index_info": {
    "dimension": 768,
    "total_vectors": 1500,
    "index_type": "IndexFlatL2",
    "created_at": "2024-01-15T11:00:00Z"
  },
  "metadata": [
    {
      "chunk_id": "chunk-001",
      "document_id": "doc-123",
      "chunk_index": 0,
      "start_char": 0,
      "end_char": 1000,
      "faiss_index": 0
    }
  ]
}
```




Comprehensive Usage Guide

Basic Workflow Example

```
import asyncio
import aiohttp

async def complete_workflow_example():
    """Complete example of crawling, embedding, and querying."""

    base_url = "http://localhost:5002"

    async with aiohttp.ClientSession() as session:

        # 1. Start crawling a documentation site
        print("🕷️ Starting crawl...")
        crawl_data = {
            "url": "https://docs.livekit.io",
            "domain_name": "livekit-docs",
            "max_depth": 5,
            "max_pages": 100
        }

        async with session.post(f"{base_url}/crawl", json=crawl_data) as resp:
            crawl_result = await resp.json()
            crawl_task_id = crawl_result["task_id"]
            print(f"Crawl task started: {crawl_task_id}")

        # 2. Monitor crawling progress
        while True:
            async with session.get(f"{base_url}/tasks/{crawl_task_id}") as resp:
                task_status = await resp.json()
                print(f"Crawl progress: {task_status['status']} - {task_status.get('progress', 0)}%")

                if task_status["status"] in ["completed", "failed"]:
                    break

            await asyncio.sleep(5)

        if task_status["status"] == "failed":
            print("❌ Crawling failed!")
            return
```

```

print(f"✅ Crawled {task_status.get('result', {}).get('total_documents', 0)} documents"

# 3. Generate embeddings
print("💡 Generating embeddings...")
embed_data = {
    "domain": "livekit-docs",
    "force_rebuild": False
}

async with session.post(f"{base_url}/embed", json=embed_data) as resp:
    embed_result = await resp.json()
    embed_task_id = embed_result["task_id"]
    print(f"Embedding task started: {embed_task_id}")

# 4. Monitor embedding progress
while True:
    async with session.get(f"{base_url}/tasks/{embed_task_id}") as resp:
        task_status = await resp.json()
        print(f"Embedding progress: {task_status['status']} - {task_status.get('progress')}%")

        if task_status["status"] in ["completed", "failed"]:
            break

    await asyncio.sleep(5)

if task_status["status"] == "failed":
    print("❌ Embedding generation failed!")
    return

print(f"✅ Generated embeddings for {task_status.get('result', {}).get('total_chunks', 0)} chunks")

# 5. Query the documentation
print("🔍 Querying documentation...")
queries = [
    "How do I set up authentication?",
    "What are the rate limits?",
    "How to handle real-time events?"
]

for query in queries:
    query_data = {
        "query": query,

```

```
        "domain": "livekit-docs",
        "top_k": 3,
        "include_context": True
    }

async with session.post(f"{base_url}/query", json=query_data) as resp:
    result = await resp.json()

    print(f"\n? Query: {query}")
    print(f"\n? Answer: {result['answer'][:200]}...")
    print(f"\n? Sources: {len(result['sources'])} documents")
    print(f"\n⚡ Response time: {result['metadata']['response_time_ms']}ms")

# Run the complete workflow
asyncio.run(complete_workflow_example())
```

Multi-Domain Query Example

```
async def multi_domain_example():
    """Example of querying across multiple domains."""

    base_url = "http://localhost:5002"

    async with aiohttp.ClientSession() as session:
        # Query multiple domains simultaneously
        query_data = {
            "query": "How to implement real-time messaging?",
            "domains": ["livekit-docs", "docs-streamlit-io", "twilio-com"],
            "top_k": 10,
            "per_domain_k": 4,
            "include_context": True,
            "merge_strategy": "score_based"
        }

        async with session.post(f"{base_url}/query/multi-domain", json=query_data) as resp:
            result = await resp.json()

            print(f"🔍 Multi-domain query results:")
            print(f"👤 Answer: {result['answer']}")
            print(f"📊 Total sources: {len(result['sources'])}")

            # Show sources by domain
            domain_counts = {}
            for source in result['sources']:
                domain = source['domain']
                domain_counts[domain] = domain_counts.get(domain, 0) + 1

            print("📊 Sources by domain:")
            for domain, count in domain_counts.items():
                print(f" - {domain}: {count} sources")

    asyncio.run(multi_domain_example())
```

Core Components Deep Dive

1. Web Crawler (`src/crawler/`)

The crawler is built on `aiohttp` for high-performance async operations:

Key Features:

- **Concurrent crawling:** Configurable concurrent request limits
- **Smart rate limiting:** Respectful crawling with delays
- **Retry logic:** Exponential backoff for failed requests
- **Content filtering:** Domain restrictions and content validation
- **Metadata extraction:** Comprehensive page metadata capture

Crawler Configuration:

```
# In .env file
MAX_CONCURRENT_REQUESTS=10      # Concurrent requests per domain
REQUEST_TIMEOUT=30                # Timeout per request (seconds)
RETRY_ATTEMPTS=3                  # Max retry attempts
DELAY_BETWEEN_REQUESTS=1          # Delay between requests (seconds)
```

2. Content Parser (`src/crawler/content_parser.py`)

Uses `trafilatura` for clean text extraction:

Features:

- **Clean text extraction:** Removes ads, navigation, footers
- **Metadata preservation:** Titles, links, publication dates
- **Link classification:** Internal vs external link identification
- **Content validation:** Quality checks and minimum content requirements

3. Storage Manager (`src/storage/`)

Domain-based storage with dual persistence:

Key Capabilities:

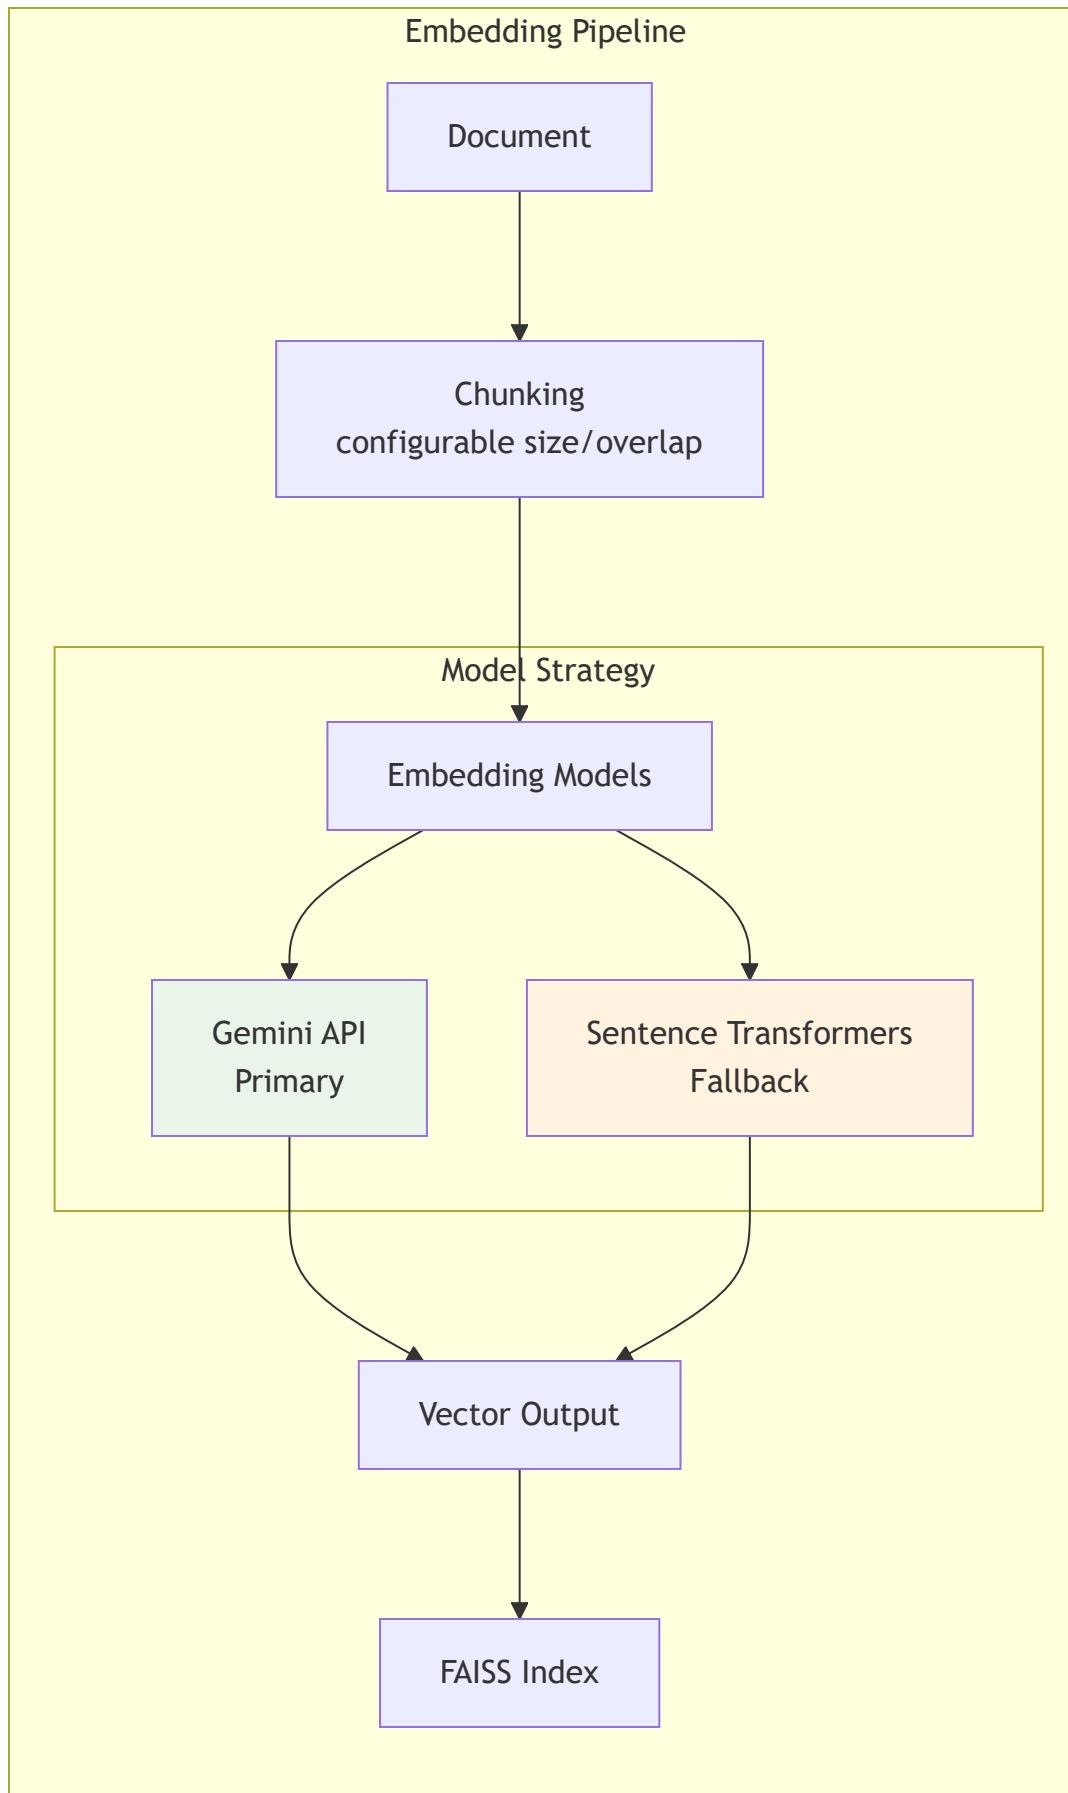
- **Automatic domain folders:** Creates organized directory structure

- **Dual format saving:** JSON for machines, YAML for humans
- **Individual file storage:** Each document as separate file
- **Efficient retrieval:** Fast document loading and querying
- **Session tracking:** Complete crawl session metadata

4. Embedding Service (`src/embeddings/`)

Multi-model embedding generation with fallback:

Architecture:



Configuration Options:

```
CHUNK_SIZE=1000          # Characters per chunk
CHUNK_OVERLAP=100         # Overlap between chunks
MAX_CHUNKS_PER_DOC=50     # Limit chunks per document
```

5. Vector Store (`src/embeddings/vector_store.py`)

FAISS-based similarity search:

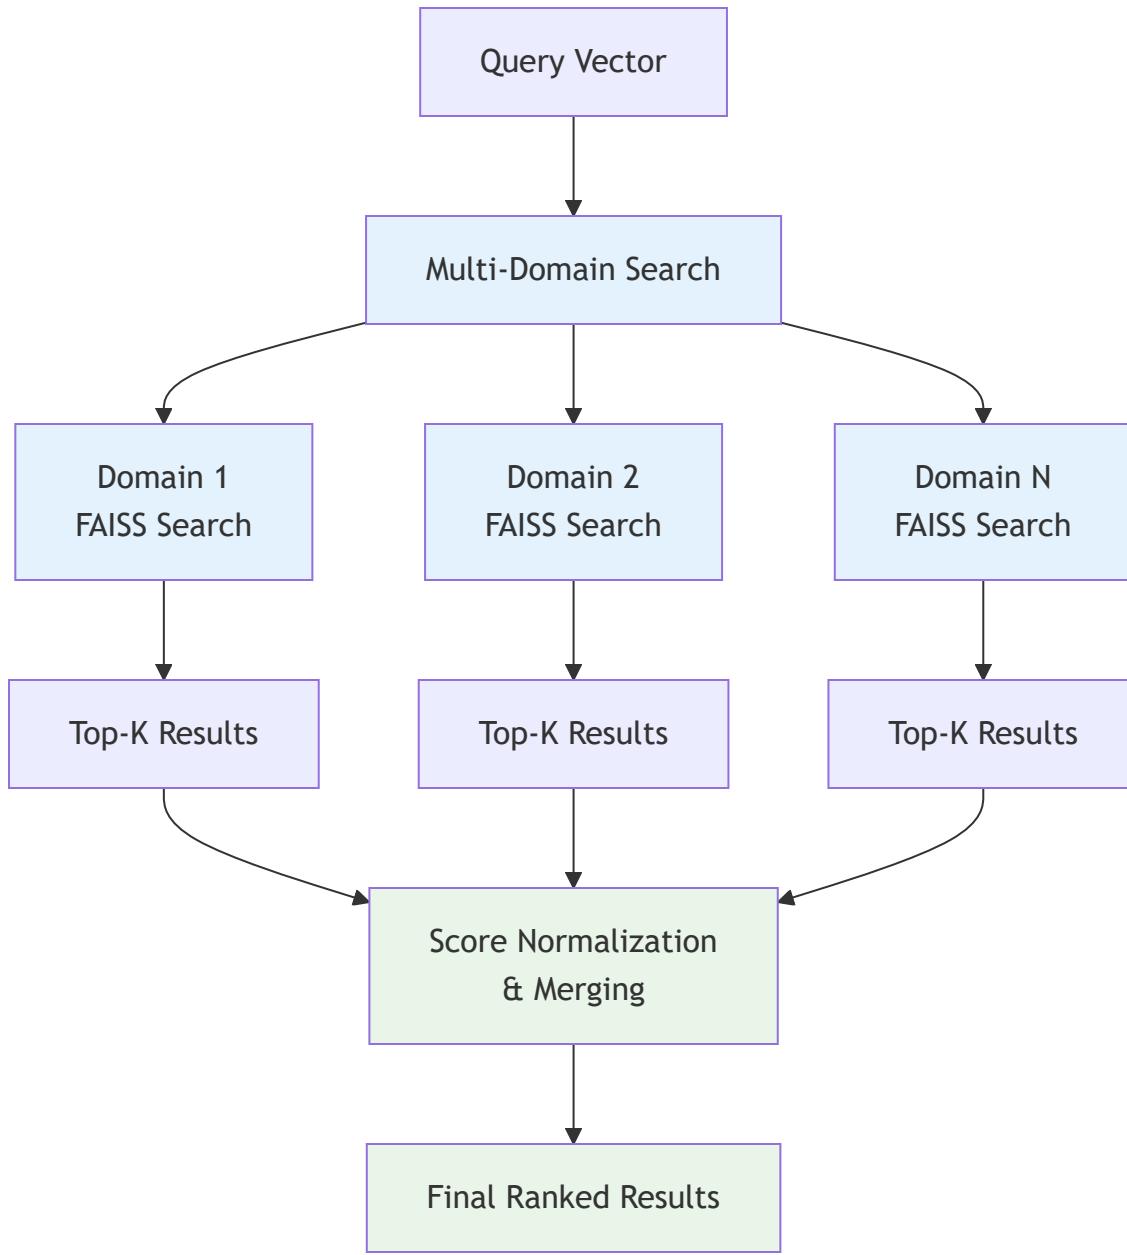
Features:

- **High-performance search:** Optimized FAISS indexes
- **Metadata integration:** Chunk-to-document mapping
- **Persistent storage:** Serialized indexes to disk
- **Score normalization:** Consistent scoring across domains

6. Multi-Domain Vector Store

Orchestrates cross-domain search:

Search Strategy:



7. RAG Pipeline (`src/qa/rag_pipeline.py`)

Complete question-answering system:

Pipeline Steps:

1. **Query embedding:** Convert question to vector
2. **Similarity search:** Find relevant document chunks
3. **Context building:** Assemble context from top results
4. **Prompt construction:** Build LLM prompt with context
5. **Answer generation:** Generate response with Gemini

6. **Source attribution:** Link answer back to sources

RAG Configuration:

```
# Query parameters
top_k=5                                # Number of chunks to retrieve
include_context=True                      # Include source context in response
min_relevance_score=0.0                  # Minimum similarity score threshold
```

Monitoring and Debugging

Logging System

Comprehensive logging across all components:

```
# Log levels available
LOG_LEVEL=DEBUG                          # Verbose debugging information
LOG_LEVEL=INFO                           # General information (default)
LOG_LEVEL=WARNING                         # Warning messages only
LOG_LEVEL=ERROR                           # Error messages only
```

Log Files Location:

```
data/logs/
├── crawler_YYYYMMDD.log      # Web crawling logs
├── embedding_YYYYMMDD.log    # Embedding generation logs
├── api_YYYYMMDD.log          # API request/response logs
└── system_YYYYMMDD.log       # General system logs
```

Background Task Monitoring

Track long-running operations:

```
# Get task status
curl http://localhost:5002/tasks/{task_id}

# Response format
{
    "task_id": "crawl-abc123",
    "status": "running",           # pending, running, completed, failed
    "progress": 45,                # Percentage complete
    "started_at": "2024-01-15T10:00:00Z",
    "completed_at": null,
    "result": null,               # Available when completed
    "error": null                 # Available if failed
}
```

System Health Monitoring

```
# System status endpoint
curl http://localhost:5002/status

# Response includes
{
    "status": "healthy",
    "uptime_seconds": 3600,
    "active_tasks": 2,
    "available_domains": ["domain1", "domain2"],
    "domains_with_embeddings": ["domain1"],
    "storage_info": {
        "total_documents": 1500,
        "total_size_mb": 45.2
    },
    "memory_usage": {
        "current_mb": 512,
        "peak_mb": 768
    }
}
```



Deployment

Render Deployment

This application is configured for deployment on Render with both Docker and native Python support.

Files Created for Deployment:

- `Procfile` - Process definition for native Python deployment
- `Dockerfile` - Container configuration for Docker deployment
- `gunicorn.conf.py` - Gunicorn configuration for production WSGI server
- Updated `requirements.txt` - Added gunicorn dependency

Deployment Steps:

1. **Connect Repository:** Connect your GitHub repository to Render
2. **Choose Deployment Method:**
 - **Docker:** Select "Docker" as runtime, uses `Dockerfile`
 - **Native Python:** Select "Python", uses `Procfile` and `requirements.txt`
3. **Environment Variables:** Set the following in Render dashboard:

`GEMINI_API_KEY=your_gemini_api_key`
`PYTHONPATH=/app/src`
`LOG_LEVEL=INFO`

4. **Build & Deploy:** Render will automatically build and deploy your application

Production Configuration:

- Uses Gunicorn with Uvicorn workers for optimal FastAPI performance
- Automatically scales workers based on CPU cores
- Configurable via `gunicorn.conf.py`
- Supports Render's dynamic PORT assignment



License

[Add your license information here]

Contributing

1. Fork the repository
2. Create a feature branch
3. Make your changes
4. Add tests if applicable
5. Submit a pull request

Support

For issues and questions:

- Create an issue on GitHub
- Check the logs in `data/logs/`